

Microservices

A Tutorial Session
JavaOne 2016

by Jason Swartz

 @swartzrock

CLASSPASS

MICROSERVICES

WHAT

WHY

a.k.a “WHEN”

HOW

What's A Microservice?

**Your Application,
Subdivided**

Recipe: Microservices

(makes 5-10 microservices)

Ingredients:

- A Monolithic API, Too Big To Handle

Instructions:

1. Find the seams in your application - by operation, business area, or team.
2. Split your API into multiple smaller API's, each on its own server.
3. Figure out how to deploy them w/o breaking everything.
4. Figure out how they'll talk to each other w/o breaking more things.
5. Figure out how complex data operations will work across multiple data stores.

Tasks

Done	Title	Assignee
Yes	Install web server	Eric
No	Add tls certs	Ronnie
No	Config ELB	Pete

Tasks



RDBMS


WARNING:
MONOLITH



```
graph TD; Tasks[Tasks] --- RDBMS[(RDBMS)]
```

The diagram illustrates a monolithic architecture. It features a rectangular box at the top labeled "Tasks" and a cylindrical database symbol at the bottom labeled "RDBMS". A vertical line connects the two, indicating a direct, integrated relationship between the application tasks and the database. A large, red, diagonal stamp with the words "WARNING:" and "MONOLITH" is overlaid on the diagram, highlighting the potential risks or challenges of this design.

Let's Split-Up
That Service




Tasks

Done	Title	Assignee
Yes	Install web server	Eric
No	Add tls certs	Ronnie
No	Config ELB	Pete

Tasks

Microservices!



Microservice Properties

1. Small & Maintainable
2. Private data
3. Generally Isolated

Microservice Nice-To-Haves

1. Small & Maintainable
2. Private data
3. Generally Isolated

Microservice Nice-To-Haves


1. Findable
2. Scalable
3. Resilient
4. Monitorable

Microservice Real Nice-To-Haves

Microservices That Speak To Each Other Via
Asynchronous Communication

Tasks

Microservices!



Let's See a
Real
Transaction

Web Client: PUT issues/123 Assignee=10

Issues Service: Read from issues db

Issues Service: GET users/10

Issues Service: POST users/notify/10

Issues Service: Write to issues db

Issues Service: Return response with result

Web Client: Read response & status then handle it

Web Client: PUT issues/123 Assignee=10

Issues Service: Read from issues db

// block for synchronous transaction

// block for synchronous transaction

Issues Service: Write to issues db

Issues Service: Return response with result

Web Client: Read response & status then handle it

Blocking Isn't
Isolated, Scalable
Or Resilient

But It Is
Easy!

Synchronous Messaging Pros

1. Easy!
2. Gratifying.
3. Testable.

Synchronous Messaging Cons

1. Prevents microservices from being isolated, scalable, and resilient.
2. Hope you like waiting.

Microservices
Should Talk
Asynchronously

Talking Asynchronously

1. Non-Blocking IO (eg RxJava, Netty)
2. Queues (eg SQS)
3. Event Streams (eg Kafka, Kinesis)
4. Actors (eg Akka)

What About The
Clients?

Responding Asynchronously

1. Synchronous REST with Async Clients
2. Http 202
3. Websockets
4. Http/2

That's **What**
Microservices Are All
About

WHAT

WHY

a.k.a “WHEN”

HOW

Monolith To Microservices

Why

Might This Be A
Good Idea?

**(Assuming You're Planning
to Split Up your Monolith)**

Monoliths - Not So Bad

1. Finding code and refactoring it is pretty easy.
2. No network lag
3. No format conversion errors
4. Deploy in one go

That's All

Thanks For Listening!

Monoliths - Not So Scalable Either

1. Hard to scale up just one feature
2. Hard to deploy just one change
3. Not suitable for more than a small team

**Organizations Are
Constrained To Produce
Designs That Match Their
Communication Structure**

Tasks




RDBMS

Max Capacity:
One Team

Tasks

Microservices!



Max Capacity:
Two Teams

Don't Split Monoliths
Across Your
Communication
Structures

Microservices Pros

(we kind of covered this already)

Microservices Challenges

1. Network Lag
2. Multiple Data Formats
3. Service Resolution
4. Troubleshooting
5. More Deployment Work
6. More Monitoring

No Problem.
Let's Get Started On
How
You'll Build Em

WHAT

WHY

a.k.a “WHEN”

HOW

Microservices In Java

**How Will You Get
Your App In
Production?**

Deployment Formats

1. Self-Running Jar's
2. VM Images
3. Docker Images

Deployment Strategies

1. Use A CI To Build / Test / Merge in CI
2. Use A CD To Deploy (Push-Button / Auto)
3. VM Deployment (eg Spinnaker)
4. Docker Deployment (eg ECS, K8s, Marathon)
5. Cluster Management (eg K8s, Mesos, Swarm)

Should I Deploy To The Cloud?

Yes

Should I Test It First?

YES OH YES

Microservice Management

1. Service Resolution (K8s, Consul)
2. Scaling
3. Visibility
4. Authentication
5. Monitoring

**Let's Talk
Microservices
Development**

Choosing A Framework

(why switch now?)

Choosing A Framework

1. Spring Boot
2. Lagom
3. Ratpack
4. Dropwizard

Choosing A Data Store

1. Good Ol' SQL (eg MySQL, Postgresql)
2. Document-type NoSQL
3. Fast Key-Value Caches

Choosing A Data Scheme

1. CRUD operations endpoints & tables
2. CRUD endpoints and immutable tables
3. Event-Sourced Data

Create-Read-
Update-Delete

Let's Talk About Immutable Tables

Tasks

Done	Title	Assignee
Yes	Install web server	Eric
No	Add tls certs	Ronnie
No	Config ELB	Pete

Issue Endpoints

GET /issues

GET /issues/{id}

POST /issues

PUT /issues/{id}

How Do These Events
Affect The Database?

1. POST /issues title='Config ELB'

id	updated	title	assignee	done
1	09-18 18:13	Config ELB	NULL	false

1. POST /issues title='Config ELB'
2. PUT /issues/1 assignee=10

id	updated	title	assignee	done
1	09-18 18:16	Config ELB	10	false

1. POST /issues title='Config ELB'
2. PUT /issues/1 assignee=10
3. PUT /issues/1 done=true

id	updated	title	assignee	done
1	09-18 18:24	Config ELB	NULL	true

1. POST /issues title='Config ELB'
2. PUT /issues/1 assignee=10
3. PUT /issues/1 done=true

Not Bad.

id	updated	title	assignee	done
1	09-18 18:24	Config ELB	NULL	true

1. `POST /issues title='Config ELB'`
2. `PUT /issues/1 assignee=10`
3. `PUT /issues/1 done=true`

Do You Know How We Got Here?

id	updated	title	assignee	done
1	09-18 18:24	Config ELB	NULL	true

Do You Know How We Got Here?

id	updated	title	assignee	done
1	09-18 18:24	Config ELB	NULL	true

1. `POST /issues title='Config ELB'`
2. `PUT /issues/1 assignee=10`
3. `PUT /issues/1 done=true`

Do You Know How We Got Here?

id	updated	title	assignee	done
1	09-18 18:24	Config ELB	NULL	true

1. POST /issues title='Config ELB'
2. PUT /issues/1 assignee=10
3. PUT /issues/1 done=true

Why is 'assignee' NULL?

id	updated	title	assignee	done
1	09-18 18:24	Config ELB	NULL	true

Mutable Table Rows Lose History

Immutable Table

Rows **Keep** Their History

Let's Try To
Lock Down
Our State

1. POST /issues title='Config ELB'

id	updated	title	assignee	done
1	09-18 18:13	Config ELB	NULL	false

1. POST /issues title='Config ELB'
2. PUT /issues/1 assignee=10

id	updated	title	assignee	done
1	09-18 18:13	Config ELB	NULL	false
1	09-18 18:16	Config ELB	10	false

1. POST /issues title='Config ELB'
2. PUT /issues/1 assignee=10
3. PUT /issues/1 done=true

id	updated	title	assignee	done
1	09-18 18:13	Config ELB	NULL	false
1	09-18 18:16	Config ELB	10	false
1	09-18 18:19	Config ELB	NULL	false
1	09-18 18:24	Config ELB	NULL	true

1. POST /issues title='Config ELB'
2. PUT /issues/1 assignee=10
3. PUT /issues/1 done=true

id	updated	title	assignee	done
1	09-18 18:13	Config ELB	NULL	false
1	09-18 18:16	Config ELB	10	false
1	09-18 18:19	Config ELB	NULL	false
1	09-18 18:24	Config ELB	NULL	true

1. GET /issues/1

id	updated	title	assignee	done
1	09-18 18:13	Config ELB	NULL	false
1	09-18 18:16	Config ELB	10	false
1	09-18 18:19	Config ELB	NULL	false
1	09-18 18:24	Config ELB	NULL	true

How Do You Make An Append-Only Table?

One: Don't Let Your DB
User Make
Changes

```
Grant select, insert on  
issues to my-db-user;
```

Two: Pick The Right Columns

```
create table issues (  
    id serial,  
    created timestamp default now(),  
    creator_id int,  
    issue_id int default nextval('iseq'),  
    title text,  
    assignee int,  
    done boolean default false  
)
```

**Do We Still Need
State?**

Let's Talk About Event-Sourcing

1. POST /issues title='Config ELB'
2. PUT /issues/1 assignee=10
3. PUT /issues/1 done=true

id	updated	title	assignee	done
1	09-18 18:13	Config ELB	NULL	false
1	09-18 18:16	Config ELB	10	false
1	09-18 18:24	Config ELB	10	true

1. POST /issues title='Config ELB'
2. PUT /issues/1 assignee=10
3. PUT /issues/1 done=true

Events

id	updated	title	assignee	done
1	09-18 18:13	Config ELB	NULL	false
1	09-18 18:16	Config ELB	10	false
1	09-18 18:24	Config ELB	10	true

States

Events

States

Create-Issue



Issue-Created

Insert

Assign-Issue



Issue-Assigned

Update

Complete-Issue



Issue-Complete

Update

Events

States

Create-Issue



Issue-Created

Insert

Assign-Issue



Issue-Assigned

Insert

Complete-Issue



Issue-Complete

Insert

Events

States

Create-Issue

Insert



Issue-Created

Assign-Issue

Insert



Issue-Assigned

Complete-Issue

Insert



Issue-Complete

Events

Virtual States

Create-Issue	Insert	→	Issue-Created
Assign-Issue	Insert	→	Issue-Assigned
Complete-Issue	Insert	→	Issue-Complete

Now We're Storing
Events,
Not States

```
create table issue_events (  
    id serial,  
    created timestamp default now(),  
    issue_id int default nextval('iseq'),  
    originator text,  
    payload text  
)
```


1. `POST /issue/1/event 'Originator: 4a48239-8a..''`
`payload='<Update val="done=true">'`

id	created	issue_id	originator	payload
14	09-18 18:50	1	4a482...	<...>

Create **Events** And
Simulate The **State**

Real Events

1. Create-Issue

Virtual States

```
Issue("Config ELB", null, false);
```

Real Events

1. ~~Create-Issue~~
2. Assign-Issue

Virtual States

```
Issue("Config ELB", 10, false);
```

Real Events

1. ~~Create-Issue~~
2. ~~Assign-Issue~~
3. Complete-Issue

Virtual States

```
Issue("Config ELB", 10, true);
```

So **Why** Use
Event-Sourcing?

Reasons For Event-Sourcing

1. High Write Performance
2. Potential for Command/Query Separation
3. Auditable
4. Replayable
5. Undo-able
6. Monitorable

It's Like Having **Control**
Over The **Versions** of
Your State Changes

It's Like Having Control
Over The Versions of
Your Data

It's Like **Git**
For Your Data

Reasons Against Event-Sourcing

1. Frankly, It's Weird
2. Requires Events. No Events, No Event-Sourcing.
3. As Of Sept 2016, It's Still Non-Standard

That About Sums Up
**Microservice
Development**

That About Sums Up
**Microservice
Development**

Okay, Actually That's The
Entire Session

Unless There's More Time

WHAT

WHY

a.k.a “WHEN”

HOW

by Jason Swartz

 @swartzrock

CLASSPASS

Microservices

A Tutorial Session
JavaOne 2016

Make Sure To See...

Web Applications for the REST of Us: An Introduction to Ember.js, Akka, and Spray (Craig Tatarzyn & Sean Kowaski)

A Practical RxJava Example with Ratpack (Laurent Doguin)

One Microservice Is No Microservice: They Come In Systems
(Markus Eisele)

**Thank You
For Attending**

Fin


Just Kidding, There's
Another Section

Let's Talk

Asynchronous

Architectures in **AWS**

Do You Want To Use
Caches To Prevent
Unnecessary
Blocking Calls?




Do You Want To Insert
Rows And Publish
Events In One Stroke?




**DynamoDB
Database**



**Kinesis
Streams**



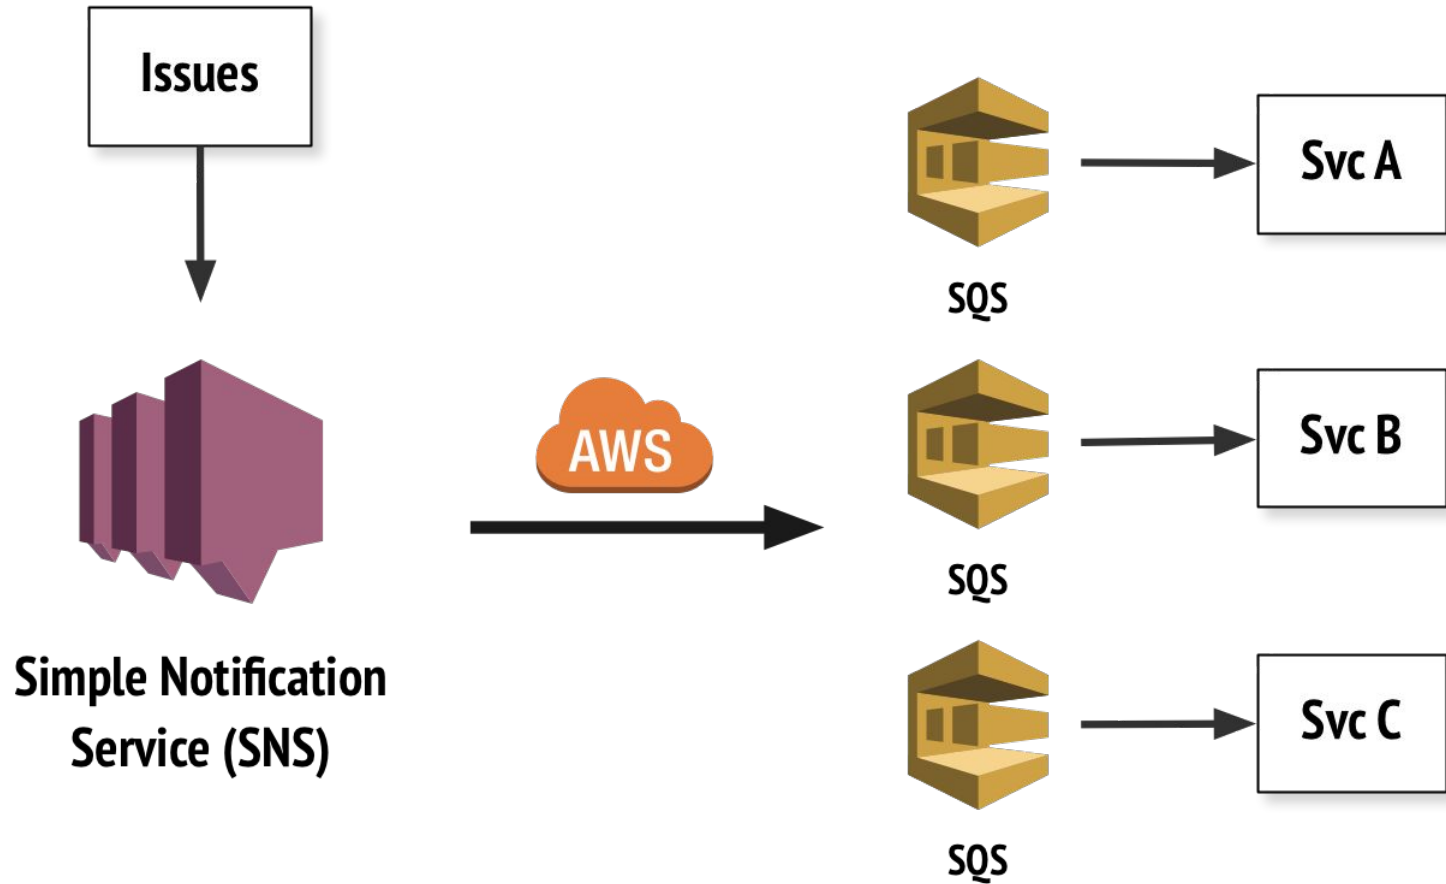
Do You Want To
Mass-Publish **One** Event
To **All** Microservices?



**Simple Notification
Service (SNS)**



**Simple Queue
Service (SQS)**



Note About Your Solution:
I'm Not Saying It's **AWS**

But It's **AWS**

Okay, We Are Really
Done This Time

**THIS SPACE
LEFT BLANK**